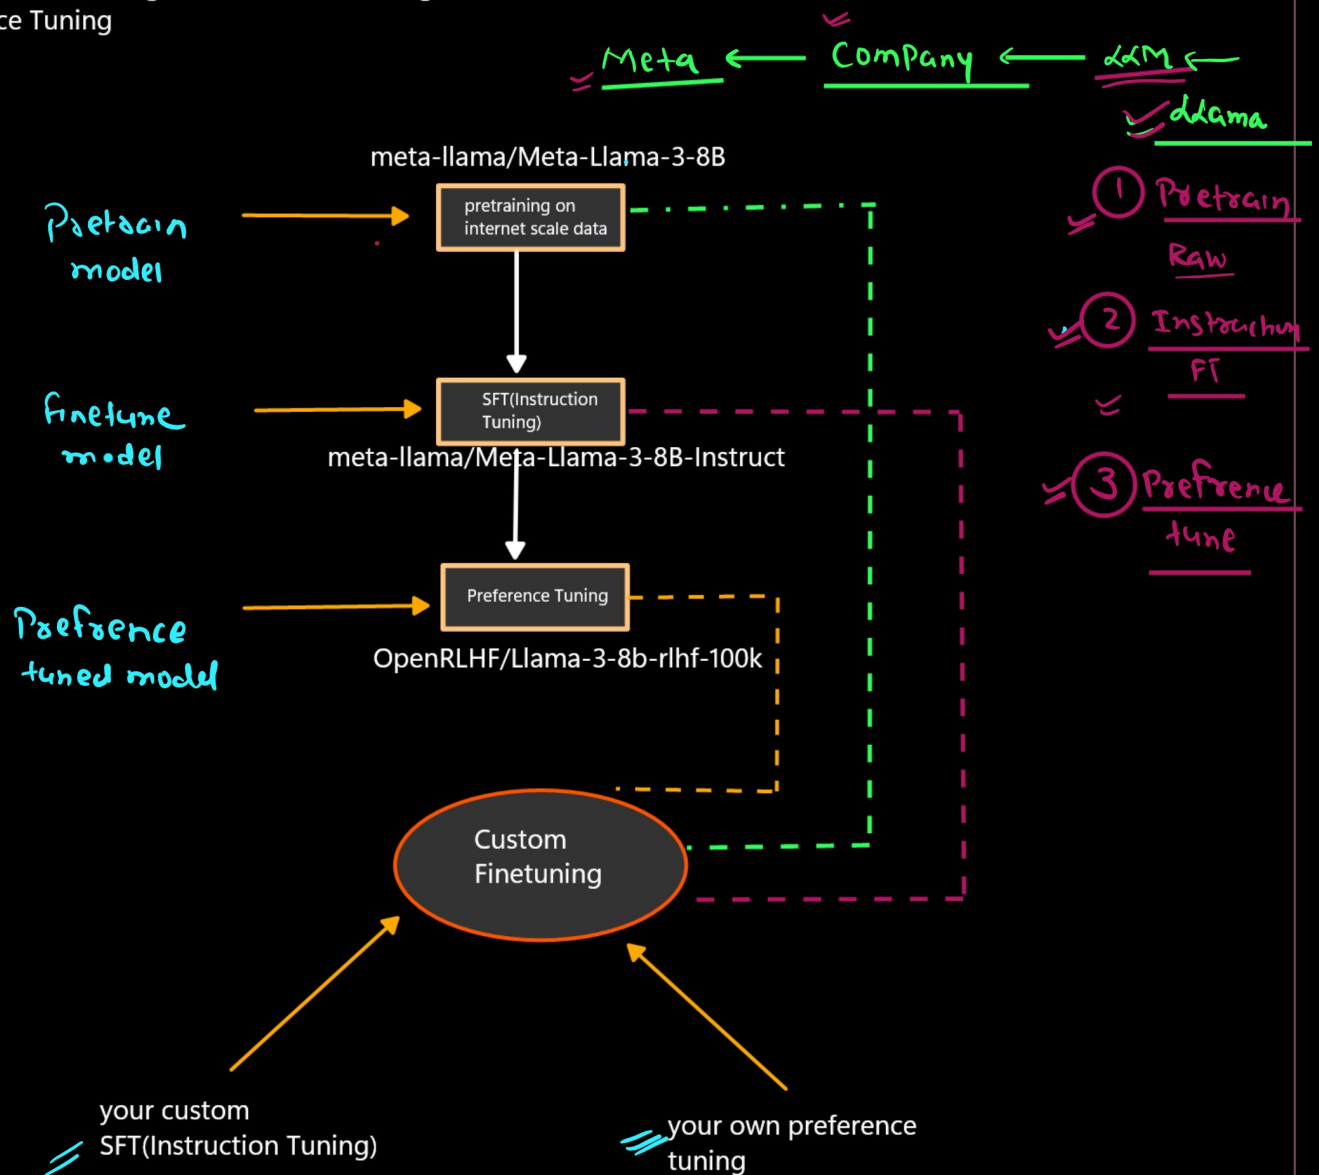


Fine-tuning

Fine-tuning is the process of taking a already trained model and training it further on domain-specific or task-specific data so that it performs better for a particular use case.



LLM Training Pipeline

1. Pretraining of model
2. Supervised Finetuning(Instruction Finetuning)
3. Preference Tuning



1. hugging face Modules
2. unsloth

1. FULL Finetuning
2. Partial Finetuning(PEFT)



meta-llama/Meta-Llama-3-8B



meta-llama/Meta-Llama-3-8B-Instruct
(how to follow the instruction)

usecase: Pharama --> document on molecule study

1. to train a llm from scratch it is difficult
2. i will take any opensource model i will download it on my owm server
3. and then i will finetune that


my requirement was just to generate the data not the conversation AI/Chatbot/QA

my requirement was just to generate the data

meta-llama/Meta-Llama-3-8B --> i have trained(finetuning) on my own pharama dataset

conversation AI/Chatbot/QA

meta-llama/Meta-Llama-3-8B-Instruct --> further i will finetuning it on my own dataset
IP--OP



① Pretraining (Unsupervised / Self supervised learning)

② SFT IFT → Supervised Finetuning (ZIP ↔ OIP)

Parameter Level

① Huge GPU Power ② Huge Infrq

① Full Finetuning ← Retrain all the Parameter (weight & bias)
② Partial Finetuning ← Subset of Parameter

Two method

① OLD School method

④ CNN → VGG / ResNet ...
BERF
BARF
TS retrain

① Freeze all the layer train only last OIP layer
② Freeze some starting layer & retrain the remaining last layer

(Latent ← LLM ← Huge architecture)

Latent

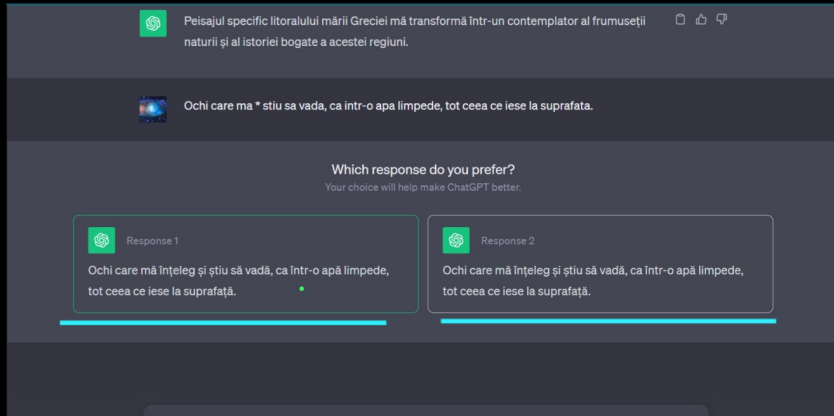
② PEFT ← Small Infrq
Parameter efficient. FT
Single GPU | Small VRAM

Low Rank Adaption ① LORA → Q LORA
↳ ADAPTER


(weight-decomposition) ② DORA Quantized
(Float → small float)
↳ int

③ BitBl
④ IA³ } X

3 Preference alignment



← collect-
Again we train
Refine our model



Hugging Face is an AI company that provides a complete ecosystem for AI development

Mainly including training LLMs from scratch, fine-tuning, evaluation, hosting applications.

almost the entire modern LLM fine-tuning industry workflow revolves around the Hugging Face ecosystem.

In modern Hugging Face ecosystem: **Most Important Libraries**

"If your goal is only modern LLM fine-tuning..."

Then these libraries are MUST KNOW:

Transformers

Datasets

Tokenizers

PEFT

Accelerate

TRL

Bitsandbytes

Safetensors

Evaluate

Lighteval

Evaluation Matrix:

accuracy

F1-score

BLEU

ROUGE

perplexity

benchmark scores

Native vs Outsider Library

"Now here is an interesting point..."

Not every tool is originally built by Hugging Face."

Native HF Libraries

Transformers

Datasets

Tokenizers

PEFT

Accelerate

TRL

Diffusers

Evaluate

Optimum

Safetensors

Outsider but Integrated

Bitsandbytes

timm

Sentence Transformers

Argilla

Distilabel

So the ecosystem is integrated, even if ownership is different.

- Datasets Library/ Custom Dataset
- Preprocessing / Cleaning
- Tokenizer (Tokenized Dataset)
- Transformers Model Loading
- PEFT (LoRA Adapters)
- Bitsandbytes (4-bit / Quantization)
- Accelerate (Distributed & Multi-GPU Training)
- Trainer / TRL (SFT, RLHF, DPO, PPO, GRPO)
- Evaluation (Accuracy, BLEU, ROUGE, Perplexity, Benchmarks)
- Save Checkpoints (Safetensors)
- Push to Hugging Face Hub
- Inference / Deployment (TGI, TEI, Inference Endpoints)
- Monitoring & Iteration

Accuracy / Quality Monitoring

Latency Monitoring

GPU / CPU Resource Monitoring

Token Usage / Cost Monitoring

Add new training data

Better LoRA adapters

RLHFDPO tuning

Hyperparameter tuning

Hugging Face is the foundation of the Finetuning ecosystem.

And Unsloth is an ultra-optimized performance layer built on top of that foundation which enables:

faster training

lower VRAM usage

longer context training

cheaper fine-tuning

Now Enters Unsloth

"Everything we discussed so far...
was the standard Hugging Face workflow.

But there was some huge issue::
slow
consumes huge VRAM
expensive

And that is exactly the problem solved by:
UNSLOTH."

Unsloth Definition

"Unsloth is an ultra-optimized LLM fine-tuning framework built on top of the Hugging Face ecosystem that enables fast LoRA/QLoRA/RLHF training with very low VRAM usage."

Main Claims of Unsloth

"Unsloth claims:

2x to 3x faster training
50–80% less VRAM usage
Ability to train large models even on free GPUs"
Real Practical Meaning

"For example:

If a Hugging Face workflow takes:

10 hours
40 GB VRAM

Then Unsloth may complete the same task in:

~5 hours
12–16 GB VRAM."
LLaMA Example

"Take the example of LLaMA 3.1 8B:

Standard HF:
24 GB VRAM

Unsloth:
only ~7–8 GB."

"The most shocking feature of Unsloth is:

LONG CONTEXT TRAINING."

20K–300K+ token training becomes possible.

"Where standard HF gives OOM...
Unsloth can handle extremely long contexts."

GPU vs Context Length

"Example:
LLaMA 3.1 (8B)

12 GB → ~21K tokens

16 GB → ~40K

24 GB → ~78K tokens

80 GB → ~340K tokens

Whereas standard HF:
Even on an 80 GB GPU, only ~28K tokens are supported."

But HOW is Unsloth So Fast?

"Now the question is...

How is this magic happening?"

Answer:

Deep low-level GPU optimizations.

CUDA and Triton

"CUDA is NVIDIA's low-level GPU programming layer.

It:

runs on GPUs

is massively parallel

is extremely fast

provides low-level control

And Triton is a Python-like GPU language that generates highly optimized CUDA kernels."

Internal Optimizations

"Internally, Unsloth uses:

custom CUDA kernels

Triton kernels

fused attention

optimized forward/backward pass

smart checkpointing

Flash Attention

manual backpropagation

automatic sequence packing"

"The most important point:

Unsloth does not rely on approximation tricks.

It maintains exact math.

That is why:

accuracy loss is negligible
model quality does not degrade"

Architecture Flow

$2.3456 \times 8.9123 = 20.90518688$

$2.3 \times 8.9 \approx 20.47$

"Exact math means Unsloth improves speed and memory efficiency without changing the actual mathematical correctness of the model computations."

"Internally, the stack looks something like this:"

Unsloth

↓

Transformers + PEFT + TRL

↓

PyTorch

↓

CUDA / Triton

↓

GPU

"This means Unsloth does not replace Hugging Face.

Instead, it works like an optimized engine on top of Hugging Face."

